

# Team 12

Reading: #12 Distributed Tracing  
Prototyping: Fog-Alert system

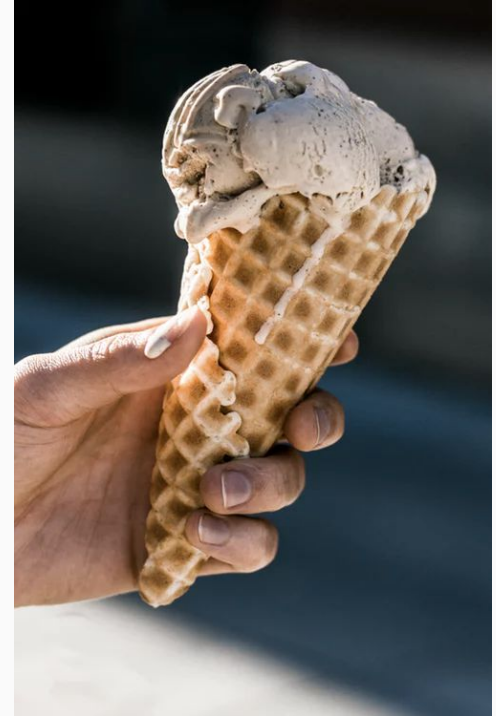
Karl Wolf, Michael Narodovitch, Pavan Chitradurga Thammanna

# Prototyping Assignment: Fog-Alert

Scenario: Managing cold storage houses in remote locations

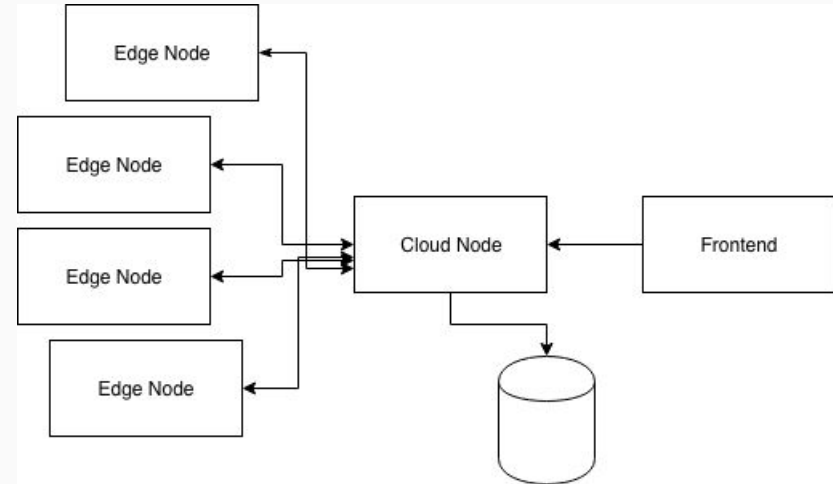
Solution: Edge devices measure temperatures

- Local Alert
- Report to Cloud Backend



# Fog Alert

- Environmental Data using Tinkerforge Sensors
- Sent live to cloud; On error: buffer & resend!
- Write to disk in case of power outage
- Chaos Penguin to make sure buffering works



# Reading Assignment: Googles Dapper

# Motivation for Tracing

## spring-boot-loader: IllegalArgumentException for file names with special characters #17208

New issue



asibross opened this issue 4 hours ago · 0 comments



asibross commented 4 hours ago



When a URL contains a special character such as `+`, the call `getRootJarFileFromUrl` fails. Here is a sample repro:

```
public static void main(String[] args) throws IOException {
    URL url = new URL("jar:file:/tmp/test/test+1.jar!/META-INF/MANIFEST.MF");
    JarFile jarFile = new Handler().getRootJarFileFromUrl(url);
}
```

This is the exception:

```
Exception in thread "main" java.io.IOException: Unable to open root Jar file 'file:/tmp/test/test+1.
    at org.springframework.boot.loader.jar.Handler.getRootJarFile(Handler.java:324)
    at org.springframework.boot.loader.jar.Handler.getRootJarFileFromUrl(Handler.java:305)
    at Main.main(Main.java:11)
Caused by: java.lang.IllegalArgumentException: File /tmp/test/test 1.jar must exist
    at org.springframework.boot.loader.data.RandomAccessDataFile$FileAccess.openIfNecessary(Rand
    at org.springframework.boot.loader.data.RandomAccessDataFile$FileAccess.<init>(RandomAccessD
    at org.springframework.boot.loader.data.RandomAccessDataFile$FileAccess.<init>(RandomAccessD
    at org.springframework.boot.loader.data.RandomAccessDataFile.<init>(RandomAccessDataFile.jav
    at org.springframework.boot.loader.jar.JarFile.<init>(JarFile.java:88)
    at org.springframework.boot.loader.jar.Handler.getRootJarFile(Handler.java:318)
    ... 2 more
```

I believe the root cause is the decoding of the root path that happens [here](#). The returned String from `decode` omits the special character.

Assignees

No one assigned

Labels

status: waiting-for-triage

Projects

None yet

Milestone

No milestone

Notifications

Subscribe



You're not receiving notifications from this thread.

2 participants



# Distributing Tracing - Introduction

Debugging of distributed applications is **hard** - especially when **we do not know**

- which services are involved
- what the services do
- what is the causality of events

Distributed tracing is a **classic problem** of (distributed) system engineering

- Google presents a **solution**: [Dapper](#)
  - Google uses Dapper for distributed tracing at production
  - Dapper is the reference design for most popular open source tracing system [Apache Zipkin](#)
- Other tracing systems: “Magpie”, “X-Trace”, “Jaeger”
- For general design considerations, see references in [Dapper](#)



# Distributing Tracing - Requirements

## Requirements for distributing tracing system

- Ubiquitous deployment
- Continuous monitoring

## Design goals for distributing tracing system

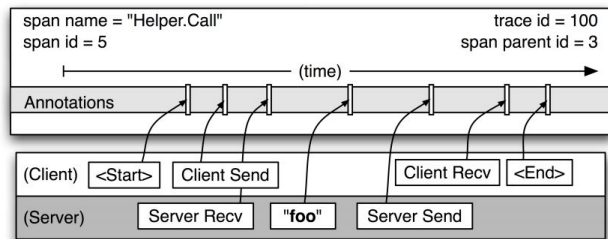
- Low overhead
- Scalability
- Low-latency-availability of tracing data
- Application-level transparency





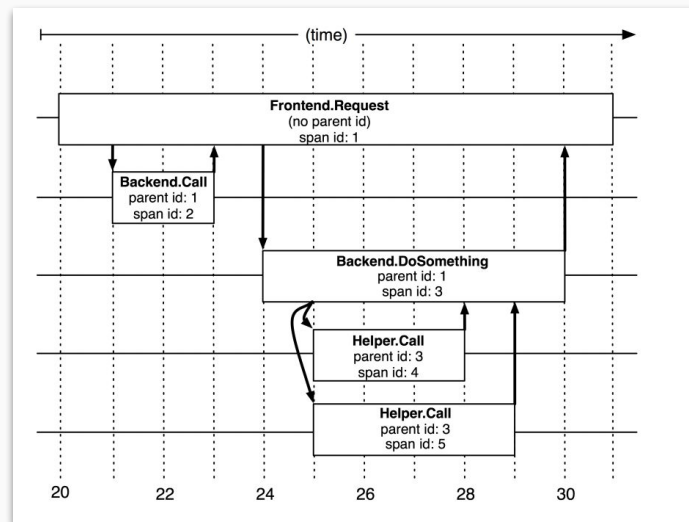
# Dapper Design - Data Structures

## Span



**Figure 3: A detailed view of a single span from Figure 2.**

## Trace Tree



- All servers communicate via GRPC
- Span is generated by GRPC instrumentation
- Annotations are added to the span by the application code (Dapper SDK...)
- Trace tree resembled from `span_id` and `parent_id`
- Globally unique ID for root spans

# Dapper Design - Trace Collection pipeline

The trace collection pipeline is characterized by a three stage process:

1. Dapper instrumentation writes to local log file
2. Dapper collectors pull log files from Dapper daemons
3. Dapper collectors push TraceTrees and Spans to BigTable repository

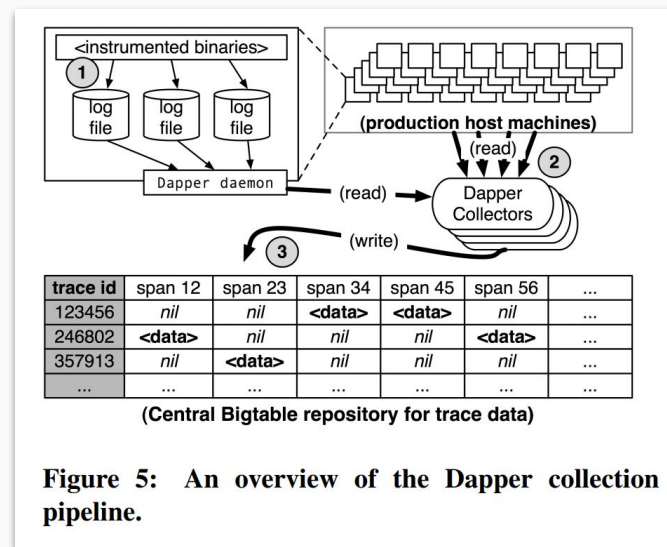


Figure 5: An overview of the Dapper collection pipeline.

# Dapper at Google Sites - Performance at Scale

## Design goals for distributing tracing system

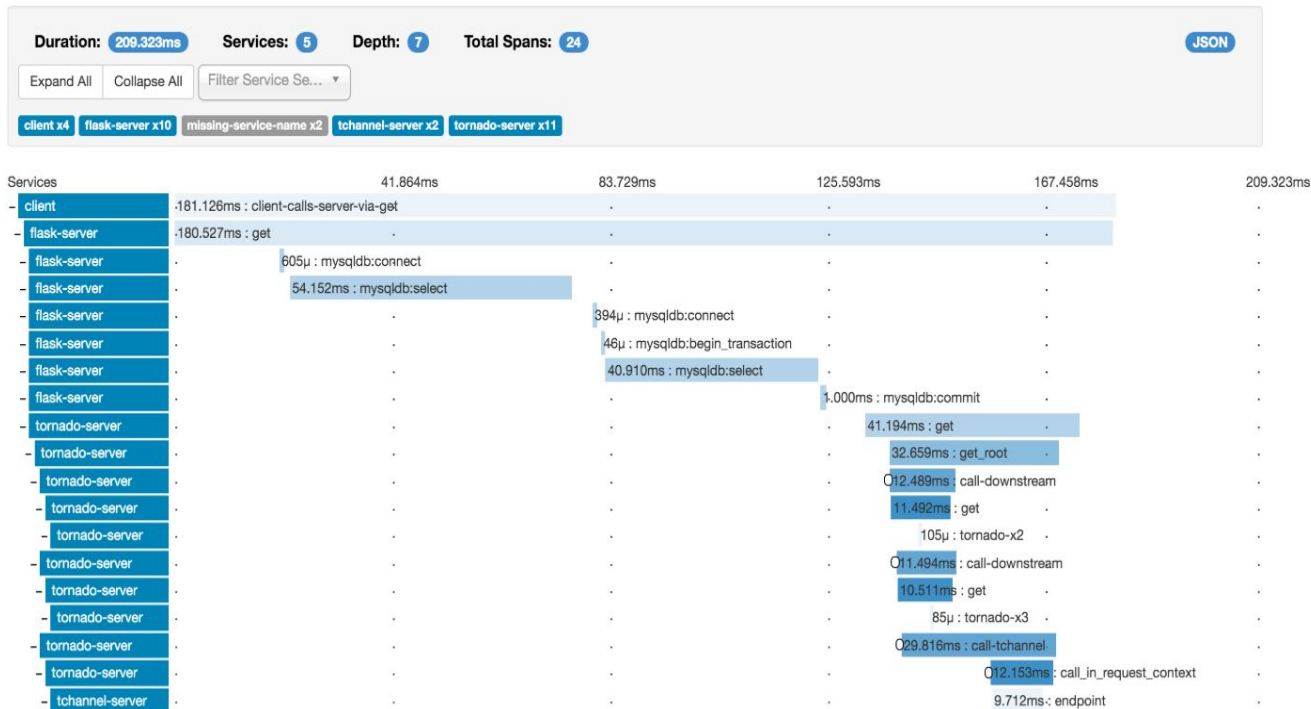
- Low overhead
- Scalability
- Low-latency-availability of tracing data
- Application-level transparency

## Reality at Google:

- Footprint of the tracing system
  - < 0.01% Network traffic, <0.3% CPU
  - Memory footprint within the noise of heap fragmentation
- Dapper Daemons deployed on >10.000 server instances
- Typical <1s end-to-end-latency for traces
- Application-level transparency by instrumentation of GRPC libraries



# Apache Zipkin



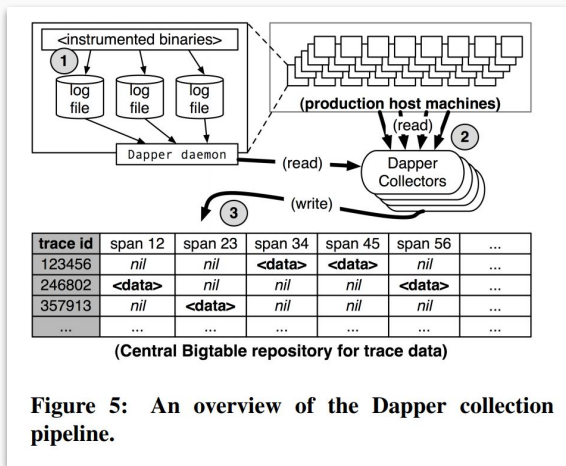
<https://zipkin.io/>

## In Zipkin Repo:

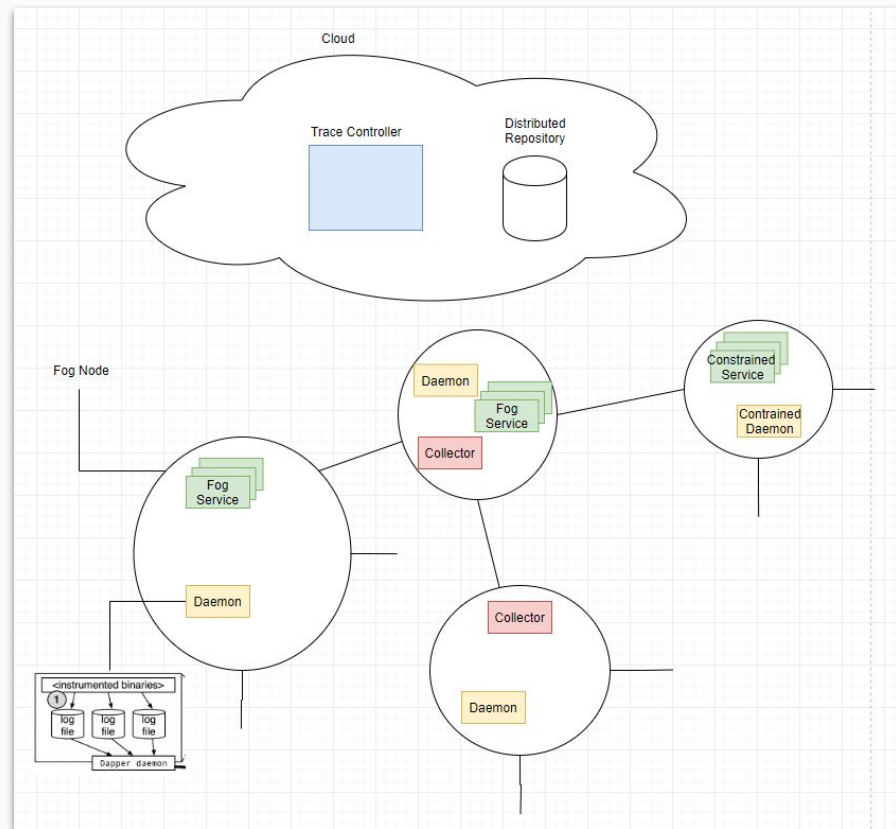
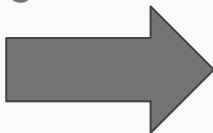
- 2029 commits
- 11,145 stars
- 237 open issues

- Based on Dapper design
- Instrumentation for (almost) every language and any transport

# Evolution to the Fog



## Fog Evolution



- Convergence of collectors and daemons to the same Fog-Node
- Location awareness collectors and daemons
- Traces still pushed to the cloud
- Trace controller for live testing